

Probabilistic Functional Programming

Donnacha Oisín Kidney

July 9, 2018

Modeling Probability

An Example

Unclear Semantics

Underpowered

Monadic Modeling

The Erwig And

Kollmansberger Approach

Other Interpreters

Theoretical Foundations

Stochastic Lambda

Calculus

Giry Monad

Other Applications

Differential Privacy

Conclusion

Modeling Probability

How do we model stochastic and probabilistic processes in programming languages?

The Boy-Girl Paradox

1. Mr. Jones has two children. The older child is a girl. What is the probability that both children are girls?
2. Mr. Smith has two children. At least one of them is a boy. What is the probability that both children are boys?

The Boy-Girl Paradox

1. Mr. Jones has two children. The older child is a girl. What is the probability that both children are girls?
2. Mr. Smith has two children. At least one of them is a boy. What is the probability that both children are boys?

Is the answer to 2 $\frac{1}{3}$ or $\frac{1}{2}$?

The Boy-Girl Paradox

1. Mr. Jones has two children. The older child is a girl. What is the probability that both children are girls?
2. Mr. Smith has two children. At least one of them is a boy. What is the probability that both children are boys?

Is the answer to 2 $\frac{1}{3}$ or $\frac{1}{2}$?

Part of the difficulty in the question is that it's ambiguous: can we use programming languages to lend some precision?

Using normal features built in to the language.

```
from random import randrange, choice
```

```
class Child:
```

```
    def __init__(self):
```

```
        self.gender = choice(["boy", "girl"])
```

```
        self.age = randrange(18)
```



```
from operator import attrgetter

def mr_jones():
    child_1 = Child()
    child_2 = Child()
    eldest = max(child_1, child_2,
                  key=attrgetter('age'))
    assert eldest.gender == 'girl'
    return [child_1, child_2]
```

```
def mr_smith():  
    child_1 = Child()  
    child_2 = Child()  
    assert child_1.gender == 'boy' or \  
           child_2.gender == 'boy'  
    return [child_1, child_2]
```

What contracts are guaranteed by probabilistic functions?

What does it mean *exactly* for a function to be probabilistic?

Why isn't the following¹ “random”?

```
int getRandomNumber()  
{  
    return 4; // chosen by fair dice roll.  
              // guaranteed to be random.  
}
```

¹Randall Munroe. *Xkcd: Random Number*. en. Title text: RFC 1149.5 specifies 4 as the standard IEEE-vetted random number. Feb. 2007. URL: <https://xkcd.com/221/> (visited on 07/06/2018).

What about this?

```
children_1 = [Child(), Child()]  
children_2 = [Child()] * 2
```

How can we describe the difference between **children_1** and **children_2**?

Underpowered

There are many more things we may want to do with probability distributions.

What about expectations?

```
def expect(predicate, process, iterations=100):
    success, tot = 0, 0
    for _ in range(iterations):
        try:
            success += predicate(process())
            tot += 1
        except AssertionError:
            pass
    return success / tot
```

The Ad-Hoc Solution

```
expect(lambda children: all(child.gender == 'girl'  
                             for child in children),  
        mr_jones)  
expect(lambda children: all(child.gender == 'boy'  
                             for child in children),  
        mr_smith)
```

Monadic Modeling

What we're approaching is a DSL, albeit an unspecified one.

Three questions for this DSL:

- Why should we implement it? What is it useful for?
- How should we implement it? How can it be made efficient?
- Can we glean any insights on the nature of probabilistic computations from the language? Are there any interesting symmetries?

The Erwig And Kollmansberger Approach

First approach²:

```
newtype Dist a = Dist { runDist :: [(a, Rational)] }
```

A distribution is a list of possible events, each tagged with a probability.

²Martin Erwig and Steve Kollmansberger. “Functional Pearls: Probabilistic Functional Programming in Haskell”. In: *Journal of Functional Programming* 16.1 (2006), pp. 21–34. ISSN: 1469-7653, 0956-7968. DOI: 10.1017/S0956796805005721. URL: <http://web.engr.oregonstate.edu/~erwig/papers/abstracts.html%5C#JFP06a> (visited on 09/29/2016).

A random integer, then, is:

```
type RandInt = Dist Int
```

This lets us encode (in the types) the difference between:

```
children_1 :: [Dist Child]
```

```
children_2 :: Dist [Child]
```

As we will use this as a DSL, we need to define the language features we used above:

```
def mr_smith():  
    child_1 = Child()  
    child_2 = Child()  
    assert child_1.gender == 'boy' or \  
           child_2.gender == 'boy'  
    return [child_1, child_2]
```

1. = (assignment)
2. **assert**
3. **return**

Assignment

This is encapsulated by the “monadic bind”:

$$(\gg=) :: \text{Dist } a \rightarrow (a \rightarrow \text{Dist } b) \rightarrow \text{Dist } b$$

When we assign to a variable in a probabilistic computation, everything that comes later is conditional on the result of that assignment. We are therefore looking for the probability of the continuation given the left-hand-side; this is encapsulated by multiplication:

$$\begin{aligned} xs \gg= f = \text{Dist } [& (y, xp \times yp) \\ & | (x, xp) \leftarrow \text{runDist } xs \\ & , (y, yp) \leftarrow \text{runDist } (f \ x)] \end{aligned}$$

Assertion

Assertion is a kind of conditioning: given a statement about an event, it either occurs or it doesn't.

```
guard :: Bool → Dist ()  
guard True  = Dist [((), 1)]  
guard False = Dist []
```

Return is the “unit” value for a distribution; the certain event, the unconditional distribution.

return :: *a* → *Dist a*

return *x* = *Dist* [(*x*, 1)]

Putting it all Together

mrSmith :: *Dist* [*Child*]

mrSmith = **do**

child1 ← *child*

child2 ← *child*

guard (*gender child1* ≡ *Boy* ∨ *gender child2* ≡ *Boy*)

return [*child1*, *child2*]

expect :: (*a* → *Rational*) → *Dist a* → *Rational*

$$\text{expect } p \text{ } xs = \frac{\text{sum } [p \ x \times xp \mid (x, xp) \leftarrow \text{runDist } xs]}{\text{sum } [xp \mid (-, xp) \leftarrow \text{runDist } xs]}$$

probOf :: (*a* → *Bool*) → *Dist a* → *Rational*

probOf *p* = *expect* ($\lambda x \rightarrow$ **if** *p* *x* **then** 1 **else** 0)

$\text{probOf } (\text{all } ((\equiv) \text{ Girl} \circ \text{gender})) \text{ mrJones} \equiv \frac{1}{2}$
 $\text{probOf } (\text{all } ((\equiv) \text{ Boy} \circ \text{gender})) \text{ mrSmith} \equiv \frac{1}{3}$

Alternative Interpreters i

Once the semantics are described, different interpreters are easy to swap in. Take Monty Hall, for example:

```
data Decision = Decision {stick  :: Bool
                           ,switch :: Bool}
montyHall :: Dist Decision
montyHall = do
  car      ← uniform [1..3]
  choice1 ← uniform [1..3]
  let left  = [door | door ← [1..3], door ≠ choice1]
  let open  = head [door | door ← left, door ≠ car]
  let choice2 = head [door | door ← left, door ≠ open]
```


```
return (Decision {stick  = car  $\equiv$  choice1  
                  , switch = car  $\equiv$  choice2})
```

While we can interpret it in the normal way to solve the problem:

$$\begin{aligned} \text{probOf stick montyHall} &\equiv \frac{1}{3} \\ \text{probOf switch montyHall} &\equiv \frac{2}{3} \end{aligned}$$

Alternative Interpreters iii

We could alternatively draw a diagram of the process:



Theoretical Foundations

Stochastic Lambda Calculus

It is possible³ to give measure-theoretic meanings to the operations described above.

$$\mathcal{M} \llbracket \text{return } x \rrbracket (A) = \begin{cases} 1, & \text{if } x \in A \\ 0, & \text{otherwise} \end{cases} \quad (1)$$

$$\mathcal{M} \llbracket d \ggg k \rrbracket (A) = \int_X \mathcal{M} \llbracket k(x) \rrbracket (A) d\mathcal{M} \llbracket d \rrbracket (x) \quad (2)$$

³Norman Ramsey and Avi Pfeffer. “Stochastic Lambda Calculus and Monads of Probability Distributions”. In: *29th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. Vol. 37. ACM, 2002, pp. 154–165. URL: <http://www.cs.tufts.edu/~nr/cs257/archive/norman-ramsey/pmonad.pdf> (visited on 09/29/2016).

The Giry Monad

Meas is the category of measurable spaces, where the morphisms are measurable maps. There is a functor **P** on **Meas**, where for some measurable space x , $\mathbf{P}(x)$ is the set of probability measures on x^4 . The monad for **P** can be defined with two natural transformations

$$\eta : A \rightarrow \mathbf{P}(A) \quad (3)$$

$$\mu : \mathbf{P}^2(A) \rightarrow \mathbf{P}(A) \quad (4)$$

Their definitions are from above:

$$\eta = \text{return} \quad (5)$$

$$\mu(P) = P \gg id \quad (6)$$

⁴Michèle Giry. "A Categorical Approach to Probability Theory". In: *Categorical Aspects of Topology and Analysis*. Ed. by A. Dold, B. Eckmann, and B. Banaschewski. Vol. 915. Berlin, Heidelberg: Springer Berlin Heidelberg, 1982, pp. 68–85. ISBN: 978-3-540-11211-2 978-3-540-39041-1. DOI: 10.1007/BFb0022272

The implementation of the Giriy monad is quite direct:

```
newtype Measure a = Measure ((a → ℝ) → ℝ)
```

Other Applications

It has been shown⁵ that the semantics of the probability monad suitable encapsulate *differential privacy*.

⁵Jason Reed and Benjamin C. Pierce. “Distance Makes the Types Grow Stronger: A Calculus for Differential Privacy”. In: *ACM Sigplan Notices*. Vol. 45. ACM, 2010, pp. 157–168. URL: <http://dl.acm.org/citation.cfm?id=1863568> (visited on 03/01/2017).

LINQ⁶ is an API which provides a monadic syntax for performing queries (sql, etc.)

PINQ⁷ extends this to provide *differentially private* queries.

⁶Don Box and Anders Hejlsberg. *LINQ: .NET Language Integrated Query*. en. Feb. 2007. URL:

<https://msdn.microsoft.com/en-us/library/bb308959.aspx> (visited on 07/09/2018).

⁷Frank McSherry. “Privacy Integrated Queries”. In: *Communications of the ACM* (Sept. 2010). URL: <https://www.microsoft.com/en-us/research/publication/privacy-integrated-queries-2/>.

Conclusion
