

Probabilistic Functional Programming

Donnacha Oisín Kidney

July 9, 2018

Modeling Probability

An Example

Unclear Semantics

Underpowered

Monadic Modeling

The Erwig And Kollmansberger Approach

Modeling Probability

How do we model stochastic and probabilistic processes in programming languages?

The Boy-Girl Paradox

1. Mr. Jones has two children. The older child is a girl. What is the probability that both children are girls?
2. Mr. Smith has two children. At least one of them is a boy. What is the probability that both children are boys?

The Boy-Girl Paradox

1. Mr. Jones has two children. The older child is a girl. What is the probability that both children are girls?
2. Mr. Smith has two children. At least one of them is a boy. What is the probability that both children are boys?

Is the answer to 2 $\frac{1}{3}$ or $\frac{1}{2}$?

The Boy-Girl Paradox

1. Mr. Jones has two children. The older child is a girl. What is the probability that both children are girls?
2. Mr. Smith has two children. At least one of them is a boy. What is the probability that both children are boys?

Is the answer to 2 $\frac{1}{3}$ or $\frac{1}{2}$?

Part of the difficulty in the question is that it's ambiguous: can we use programming languages to lend some precision?

Using normal features built in to the language.

```
from random import randrange, choice
```

```
class Child:
```

```
    def __init__(self):
```

```
        self.gender = choice(["boy", "girl"])
```

```
        self.age = randrange(18)
```



```
from operator import attrgetter

def mr_jones():
    child_1 = Child()
    child_2 = Child()
    eldest = max(child_1, child_2,
                  key=attrgetter('age'))
    assert eldest.gender == 'girl'
    return [child_1, child_2]
```

```
def mr_smith():  
    child_1 = Child()  
    child_2 = Child()  
    assert child_1.gender == 'boy' or \  
           child_2.gender == 'boy'  
    return [child_1, child_2]
```

Unclear semantics

What contracts are guaranteed by probabilistic functions?

What does it mean *exactly* for a function to be probabilistic?

Why isn't the following¹ “random”?

```
int getRandomNumber()  
{  
    return 4; // chosen by fair dice roll.  
              // guaranteed to be random.  
}
```

¹Randall Munroe. *Xkcd: Random Number*. en. Title text: RFC 1149.5 specifies 4 as the standard IEEE-vetted random number. Feb. 2007. URL: <https://xkcd.com/221/> (visited on 07/06/2018).

What about this?

```
children_1 = [Child(), Child()]  
children_2 = [Child()] * 2
```

How can we describe the difference between **children_1** and **children_2**?

Underpowered

There are many more things we may want to do with probability distributions.

What about expectations?

```
def expect(predicate, process, iterations=100):  
    success, tot = 0, 0  
    for _ in range(iterations):  
        try:  
            success += predicate(process())  
            tot += 1  
        except AssertionError:  
            pass  
    return success / tot
```

The Ad-Hoc Solution

```
expect(lambda children: all(child.gender == 'girl'
                             for child in children),
        mr_jones)
expect(lambda children: all(child.gender == 'boy'
                             for child in children),
        mr_smith)
```

Monadic Modeling

What we're approaching is a DSL, albeit an unspecified one.

Three questions for this DSL:

- Why should we implement it? What is it useful for?
- How should we implement it? How can it be made efficient?
- Can we glean any insights on the nature of probabilistic computations from the language? Are there any interesting symmetries?

The Erwig And Kollmansberger Approach

First approach²:

```
newtype Dist a = Dist { runDist :: [(a, Rational)] }
```

A distribution is a list of possible events, each tagged with a probability.

²Martin Erwig and Steve Kollmansberger. “Functional Pearls: Probabilistic Functional Programming in Haskell”. In: *Journal of Functional Programming* 16.1 (2006), pp. 21–34. ISSN: 1469-7653, 0956-7968. DOI: 10.1017/S0956796805005721. URL: <http://web.engr.oregonstate.edu/~erwig/papers/abstracts.html%5C#JFP06a> (visited on 09/29/2016).

A random integer, then, is:

```
type RandInt = Dist Int
```

This lets us encode (in the types) the difference between:

```
children_1 :: [Dist Child]
```

```
children_2 :: Dist [Child]
```

As we will use this as a DSL, we need to define the language features we used above:

```
def mr_smith():  
    child_1 = Child()  
    child_2 = Child()  
    assert child_1.gender == 'boy' or \  
           child_2.gender == 'boy'  
    return [child_1, child_2]
```

1. = (assignment)
2. **assert**
3. **return**

Assignment

This is encapsulated by the “monadic bind”:

$$(\gg=) :: \text{Dist } a \rightarrow (a \rightarrow \text{Dist } b) \rightarrow \text{Dist } b$$

When we assign to a variable in a probabilistic computation, everything that comes later is conditional on the result of that assignment. We are therefore looking for the probability of the continuation given the left-hand-side; this is encapsulated by multiplication:

$$\begin{aligned} xs \gg= f = \text{Dist } [& (y, xp \times yp) \\ & | (x, xp) \leftarrow \text{runDist } xs \\ & , (y, yp) \leftarrow \text{runDist } (f \ x)] \end{aligned}$$

Assertion is a kind of conditioning: given a statement about an event, it either occurs or it doesn't.

```
guard :: Bool → Dist ()  
guard True  = Dist [((), 1)]  
guard False = Dist []
```

Return is the “unit” value for a distribution; the certain event, the unconditional distribution.

return :: *a* → *Dist a*

return *x* = *Dist* [(*x*, 1)]

Putting it all Together

mrSmith :: *Dist* [*Child*]

mrSmith = **do**

child1 ← *child*

child2 ← *child*

guard (*gender child1* ≡ *Boy* ∨ *gender child2* ≡ *Boy*)

return [*child1*, *child2*]


expect :: (*a* → *Rational*) → *Dist a* → *Rational*

$$\text{expect } p \text{ } xs = \frac{\text{sum } [p \ x \times xp \mid (x, xp) \leftarrow \text{runDist } xs]}{\text{sum } [xp \mid (-, xp) \leftarrow \text{runDist } xs]}$$

probOf :: (*a* → *Bool*) → *Dist a* → *Rational*

probOf *p* = *expect* ($\lambda x \rightarrow$ **if** *p* *x* **then** 1 **else** 0)

$\text{probOf } (\text{all } ((\equiv) \text{ Girl} \circ \text{gender})) \text{ mrJones} \equiv \frac{1}{2}$
 $\text{probOf } (\text{all } ((\equiv) \text{ Boy} \circ \text{gender})) \text{ mrSmith} \equiv \frac{1}{3}$



Theoretical Basis